

*Operating
Systems:
Internals
and Design
Principles*

Chapter 6 Concurrency: Deadlock and Starvation


Ninth Edition
By William Stallings

Let's look at an example


- Suppose there is a narrow staircase.
- Two people, one going down and one going up meets up in the middle.
- Neither can move anywhere except go back.
- Each of them are waiting for another person to move.
 - Deadlock
- Suppose you have slightly wider staircase
 - What happens now?
 - If they try to move to opposite direction at the same time?
 - Livelock
 - Two or more processes or threads continuously change their states in response to each other but make no actual progress.

Deadlock

- The *permanent* blocking of a set of processes that either compete for system resources or communicate with each other
- A set of processes is deadlocked when each process in the set is blocked awaiting an event that can only be triggered by another blocked process in the set
- Permanent because none of the events is ever triggered
- No efficient solution in the general case



(a) Deadlock possible




(b) Deadlock

Figure 6.1 Illustration of Deadlock


Let's look at another example

- Suppose we have processes P, Q



P

- Get A
- Get B
- Do work
- Release A
- Release B



Q

- Get B
- Get A
- Do work
- Release B
- Release A

- Joint resource diagram




Figure 6.2 Example of Deadlock




Figure 6.3 Example of No Deadlock


Fatal Regions

- Deadlock only occurs if two processes (or more) creates a path that enters the fatal region.
- What if there are more than 2 processes that shares resources
 - We need joint progress diagrams that uses more than two dimensions.

Resource Categories

Reusable

- Can be safely used by only one process at a time and is not depleted by that use
- Processors, I/O channels, main and secondary memory, devices, and data structures such as files, databases, and semaphores



Consumable

- One that can be created (produced) and destroyed (consumed)
 - Interrupts, signals, messages, and information
 - In I/O buffers

Memory Request example

- Space is available for allocation of 200Kbytes, and the following sequence of events occur:

P1

...

Request 80 Kbytes;

...

Request 60 Kbytes;

P2

...

Request 70 Kbytes;

...

Request 80 Kbytes;

- Deadlock occurs if both processes progress to their second request

Consumable resources

- Can you produce a scenario where attempting access consumable resources by multiple processes ends up in a deadlock?

Consumable Resources Deadlock

- Consider a pair of processes, in which each process attempts to receive a message from the other process and then send a message to the other process:

P1	P2
...	...
Receive (P2);	Receive (P1);
...	...
Send (P2, M1);	Send (P1, M2);

- Deadlock occurs if the Receive is blocking

Resource categories

- Deadlocks with Reusable Resources:
common and straightforward.
- Deadlocks with Consumable Resources:
 - Less typical but possible
 - If the production and consumption of resources are interdependent.
 - Such deadlocks might involve indirect circular waits.

Conditions for Deadlock

Mutual Exclusion

- Only one process may use a resource at a time
- No process may access a resource until that has been allocated to another process

Hold-and-Wait

- A process may hold allocated resources while awaiting assignment of other resources

No Pre-emption

- No resource can be forcibly removed from a process holding it

Circular Wait


- A closed chain of processes exists, such that each process holds at least one resource needed by the next process in the chain

How to model a deadlock?


- Resource allocation graphs
- Directed graph that depicts a state of the system of resources and processes
- Each process and each resource represented by a node
- Edge directed from a process to a resource indicates a resource that has been requested by the process but not yet granted.
- Edge directed from a reusable resource node dot to a process indicates a request that has been granted




(a) Resource is requested



(b) Resource is held



(c) Circular wait



(d) No deadlock

Figure 6.5 Examples of Resource Allocation Graphs




Figure 6.6 Resource Allocation Graph for Figure 6.1b

Deadlock Approaches

- Can you think about an effective way to deal with deadlocks?

- **Deadlock avoidance**

- Do not grant a resource request if this allocation might lead to deadlock

- **Deadlock prevention**

- Disallow one of the three necessary conditions for deadlock occurrence, or prevent circular wait condition from happening

- **Deadlock detection**

- Grant resource requests, when possible, but periodically check for the presence of deadlock and take action to recover

Difference between deadlock prevention and avoidance

- **Deadlock prevention** involves designing the system and its resource allocation protocols to ensure that at least one of the necessary conditions for deadlock cannot occur.
 - Using a resource allocation graph with specific rules to avoid circular waits can help prevent deadlocks
- **Deadlock avoidance** involves dynamically examining the resource allocation state of the system to ensure that each resource request can be granted without leading to a deadlock.
 - When a process requests a resource, the system temporarily pretends to allocate the resource and checks if this allocation keeps the system in a safe state.

Deadlock Prevention

- Design a system in such a way that the possibility of deadlock is excluded
- We can remove one of the conditions of deadlock from the system.
- Two main methods:
 - Indirect
 - Prevent the occurrence of one of the three necessary conditions
 - Direct
 - Prevent the occurrence of a circular wait

Deadlock Condition Prevention

■ Mutual exclusion

- If access to a resource requires mutual exclusion, then mutual exclusion must be supported by the OS
- Some resources, such as files, may allow multiple accesses for reads but only exclusive access for writes
- Even in this case, deadlock can occur if more than one process requires write permission

■ Hold and wait

- Can be prevented by requiring that a process request all of its required resources at one time and blocking the process until all requests can be granted simultaneously

Deadlock Condition Prevention

■ No Preemption

- If a process holding certain resources is denied a further request, that process must release its original resources and request them again
- OS may preempt the second process and require it to release its resources

■ Circular Wait

- The circular wait condition can be prevented by defining a linear ordering of resource types

Deadlock Condition Prevention

- Most of these prevention strategies are inefficient or unfavorable.

Deadlock Avoidance

- Allows the three necessary conditions but makes judicious choices to assure that the deadlock point is never reached
- A decision is made dynamically whether the current resource allocation request will, if granted, potentially lead to a deadlock
- Allows the three necessary conditions but makes judicious choices to assure that the deadlock point is never reached
- Requires knowledge of future process requests

Two Approaches to Deadlock Avoidance

Deadlock Avoidance

Resource Allocation Denial

- Do not grant an incremental resource request to a process if this allocation might lead to deadlock

Process Initiation Denial

- Do not start a process if its demands might lead to deadlock

Resource Allocation Denial

- Referred to as the *banker's algorithm*
- ***State*** of the system reflects the current allocation of resources to processes
- ***Safe state*** is one in which there is at least one sequence of resource allocations to processes that does not result in a deadlock
- ***Unsafe state*** is a state that is not safe

	R1	R2	R3
P1	3	2	2
P2	6	1	3
P3	3	1	4
P4	4	2	2

Claim matrix **C**

	R1	R2	R3
P1	1	0	0
P2	6	1	2
P3	2	1	1
P4	0	0	2

Allocation matrix **A**

	R1	R2	R3
P1	2	2	2
P2	0	0	1
P3	1	0	3
P4	4	2	0

C – A

R1	R2	R3
9	3	6

Resource vector **R**

R1	R2	R3
0	1	1

Available vector **V**

(a) Initial state

Figure 6.7 Determination of a Safe State

	R1	R2	R3
P1	3	2	2
P2	0	0	0
P3	3	1	4
P4	4	2	2

Claim matrix **C**

	R1	R2	R3
P1	1	0	0
P2	0	0	0
P3	2	1	1
P4	0	0	2

Allocation matrix **A**

	R1	R2	R3
P1	2	2	2
P2	0	0	0
P3	1	0	3
P4	4	2	0

C – A

	R1	R2	R3
P1	9	3	6

Resource vector **R**

	R1	R2	R3
P1	6	2	3

Available vector **V**

(b) P2 runs to completion

Figure 6.7 Determination of a Safe State

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	3	1	4
P4	4	2	2

Claim matrix **C**

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	2	1	1
P4	0	0	2

Allocation matrix **A**

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	1	0	3
P4	4	2	0

C - A

R1	R2	R3
9	3	6

Resource vector **R**

R1	R2	R3
7	2	3

Available vector **V**

(c) P1 runs to completion

Figure 6.7 Determination of a Safe State

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	0	0	0
P4	4	2	2

Claim matrix **C**

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	0	0	0
P4	0	0	2

Allocation matrix **A**

	R1	R2	R3
P1	0	0	0
P2	0	0	0
P3	0	0	0
P4	4	2	0

C - A

R1	R2	R3
9	3	6

Resource vector **R**

R1	R2	R3
9	3	4

Available vector **V**

(d) P3 runs to completion

Figure 6.7 Determination of a Safe State

	R1	R2	R3		R1	R2	R3		R1	R2	R3
P1	3	2	2	P1	1	0	0	P1	2	2	2
P2	6	1	3	P2	5	1	1	P2	1	0	2
P3	3	1	4	P3	2	1	1	P3	1	0	3
P4	4	2	2	P4	0	0	2	P4	4	2	0
Claim matrix C				Allocation matrix A				C - A			

R1	R2	R3	R1	R2	R3
9	3	6	1	1	2
Resource vector R			Available vector V		

(a) Initial state

	R1	R2	R3		R1	R2	R3		R1	R2	R3
P1	3	2	2	P1	2	0	1	P1	1	2	1
P2	6	1	3	P2	5	1	1	P2	1	0	2
P3	3	1	4	P3	2	1	1	P3	1	0	3
P4	4	2	2	P4	0	0	2	P4	4	2	0
Claim matrix C				Allocation matrix A				C - A			

R1	R2	R3	R1	R2	R3
9	3	6	0	1	1
Resource vector R			Available vector V		

(b) P1 requests one unit each of R1 and R3

Figure 6.8 Determination of an Unsafe State

```

struct state {
    int resource[m];
    int available[m];
    int claim[n][m];
    int alloc[n][m];
}

```

(a) global data structures

```

if (alloc [i,*] + request [*] > claim [i,*])
    < error >; /* total request > claim*/
else if (request [*] > available [*])
    < suspend process >;
else { /* simulate alloc */
    < define newstate by:
    alloc [i,*] = alloc [i,*] + request [*];
    available [*] = available [*] - request [*] >;
}
if (safe (newstate))
    < carry out allocation >;
else {
    < restore original state >;
    < suspend process >;
}

```

(b) resource allocation algorithm

```

boolean safe (state S) {
    int currentavail[m];
    process rest[<number of processes>];
    currentavail = available;
    rest = {all processes};
    possible = true;
    while (possible) {
        <find a process Pk in rest such that
        claim [k,*] - alloc [k,*] <= currentavail;>
        if (found) { /* simulate execution of Pk */
            currentavail = currentavail + alloc [k,*];
            rest = rest - {Pk};
        }
        else possible = false;
    }
    return (rest == null);
}

```


(c) test for safety algorithm (banker's algorithm)

Figure 6.9 Deadlock Avoidance Logic

Deadlock Avoidance Advantages

- It is not necessary to preempt and rollback processes, as in deadlock detection
- It is less restrictive than deadlock prevention

Deadlock Avoidance Restrictions

- 
- Maximum resource requirement for each process must be stated in advance
 - Processes under consideration must be independent and with no synchronization requirements
 - There must be a fixed number of resources to allocate
 - No process may exit while holding resources

Deadlock Strategies

Deadlock prevention strategies are very conservative

- Limit access to resources by imposing restrictions on processes

Deadlock detection strategies do the opposite

- Resource requests are granted whenever possible

Deadlock Detection Algorithm

- A check for deadlock can be made as frequently as each resource request or, less frequently, depending on how likely it is for a deadlock to occur
- CES algorithm



Advantages:

- It leads to early detection
- The algorithm is relatively simple



Disadvantage

- Frequent checks consume considerable processor time

	R1	R2	R3	R4	R5
P1	0	1	0	0	1
P2	0	0	1	0	1
P3	0	0	0	0	1
P4	1	0	1	0	1

Request matrix Q

	R1	R2	R3	R4	R5
P1	1	0	1	1	0
P2	1	1	0	0	0
P3	0	0	0	1	0
P4	0	0	0	0	0

Allocation matrix A

R1	R2	R3	R4	R5
2	1	1	2	1

Resource vector

R1	R2	R3	R4	R5
0	0	0	0	1

Allocation vector

Figure 6.10 Example for Deadlock Detection

Recovery Strategies

- Abort all deadlocked processes
- Back up each deadlocked process to some previously defined checkpoint and restart all processes
- Successively abort deadlocked processes until deadlock no longer exists
- Successively preempt resources until deadlock no longer exists

Integrated Deadlock Strategy

- Rather than attempting to design an OS facility that employs only one of these strategies, it might be more efficient to use different strategies in different situations
 - Group resources into a number of different resource classes
 - Use the linear ordering strategy defined previously for the prevention of circular wait to prevent deadlocks between resource classes
 - Within a resource class, use the algorithm that is most appropriate for that class
- Classes of resources
 - Swappable space
 - Blocks of memory on secondary storage for use in swapping processes
 - Process resources
 - Assignable devices, such as tape drives, and files
 - Main memory
 - Assignable to processes in pages or segments
 - Internal resources
 - Such as I/O channels

Class Strategies

- Within each class the following strategies could be used:
 - **Swappable space**
 - Prevention of deadlocks by requiring that all of the required resources that may be used be allocated at one time, as in the hold-and-wait prevention strategy
 - This strategy is reasonable if the maximum storage requirements are known
 - **Process resources**
 - Avoidance will often be effective in this category, because it is reasonable to expect processes to declare ahead of time the resources that they will require in this class
 - Prevention by means of resource ordering within this class is also possible
 - **Main memory**
 - Prevention by preemption appears to be the most appropriate strategy for main memory
 - When a process is preempted, it is simply swapped to secondary memory, freeing space to resolve the deadlock
 - **Internal resources**
 - Prevention by means of resource ordering can be used